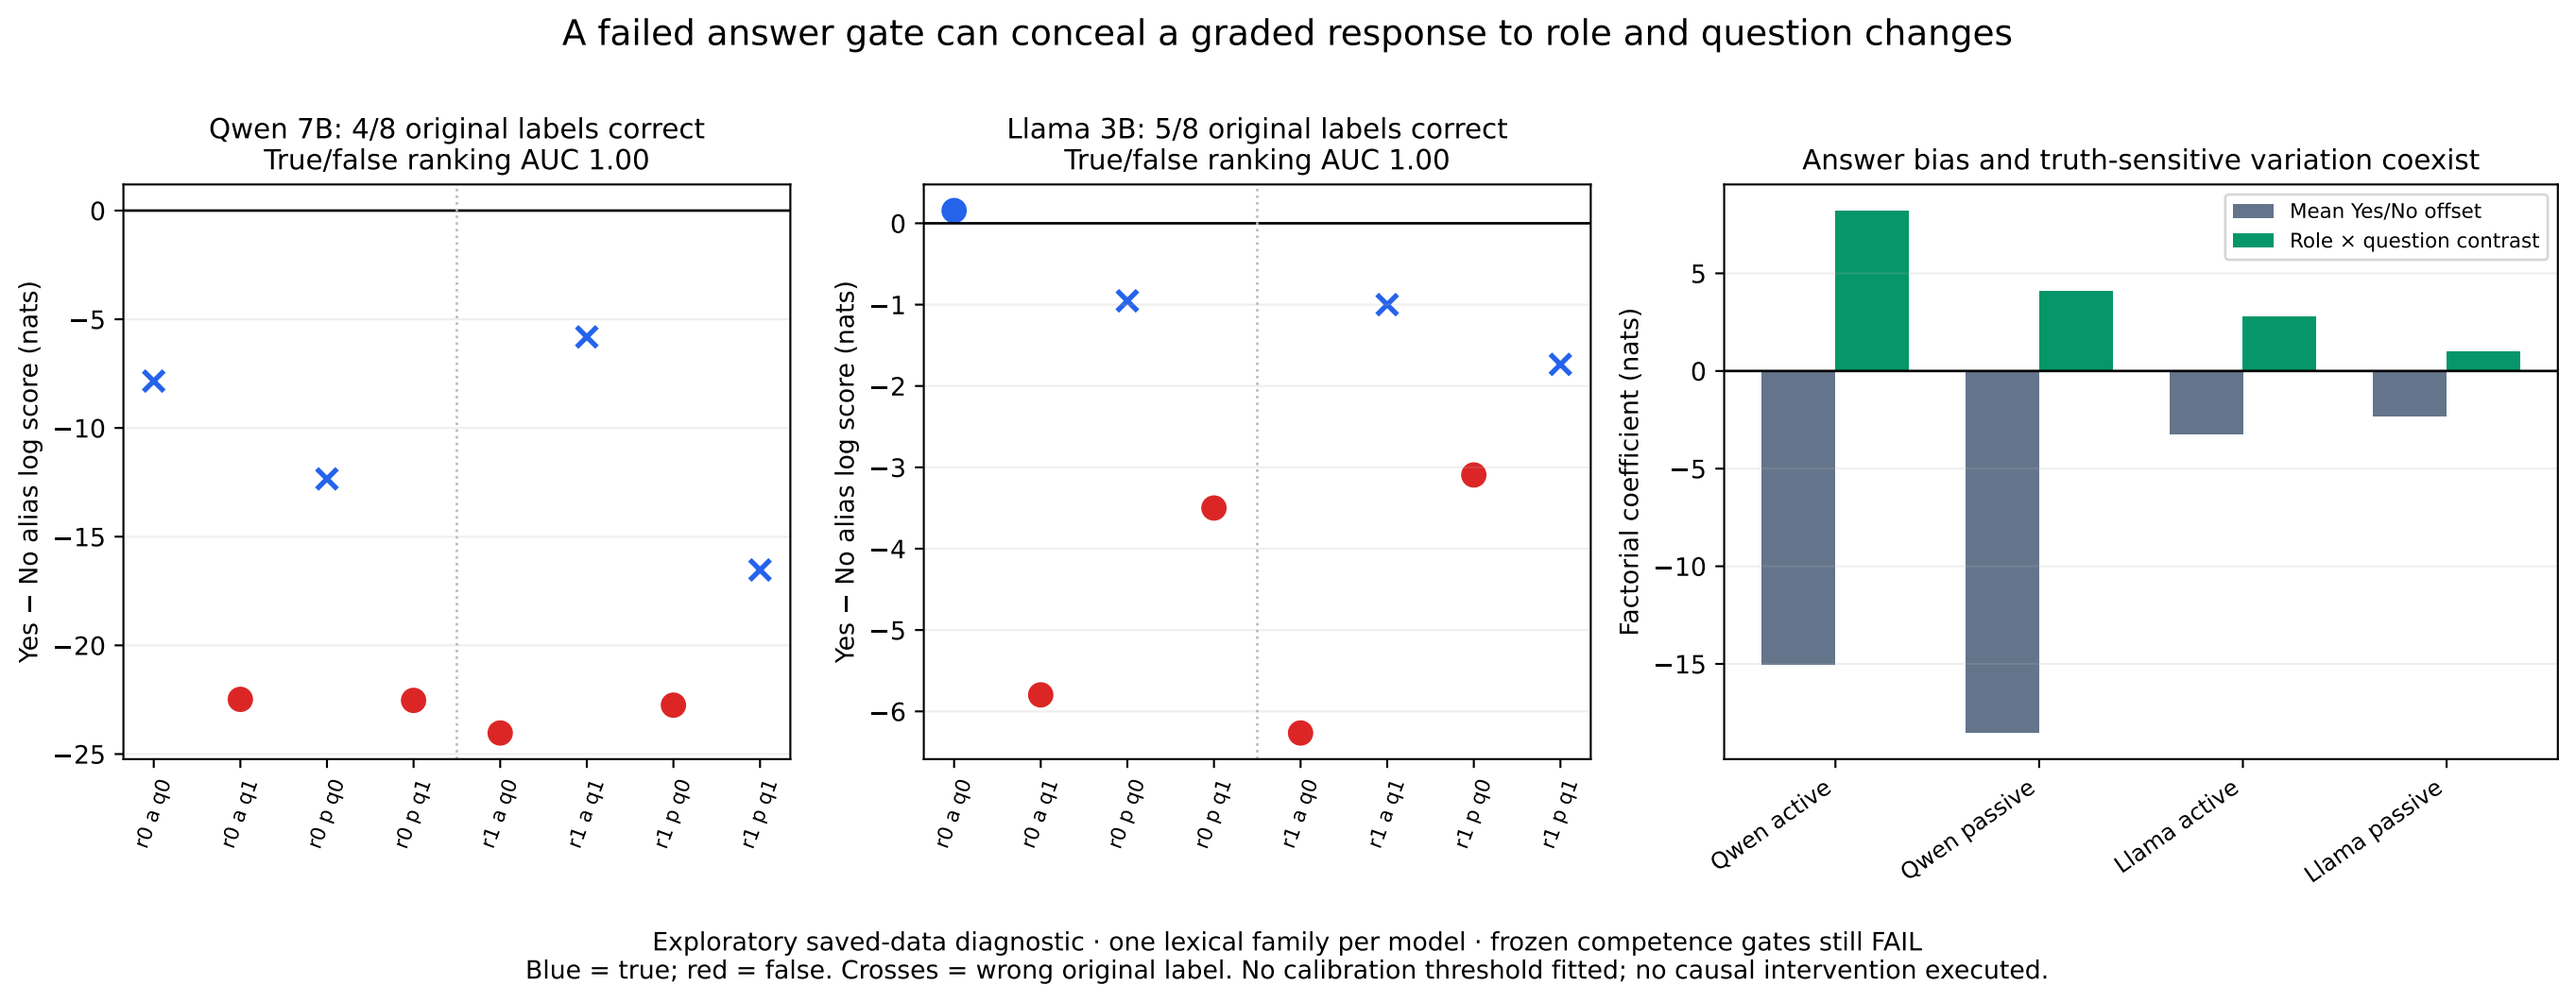
A failed answer gate can conceal a graded response to role and question changes
Qwen 7B: 4/8 original labels correct
True/false ranking AUC 1.00
Llama 3B: 5/8 original labels correct
True/false ranking AUC 1.00
Answer bias and truth-sensitive variation coexist
Yes − No alias log score (nats)
Factorial coefficient (nats)
Mean Yes/No offset
Role × question contrast
r0 a q0
r0 a q1
r0 p q0
r0 p q1
r1 a q0
r1 a q1
r1 p q0
r1 p q1
Qwen active
Qwen passive
Llama active
Llama passive
Exploratory saved-data diagnostic · one lexical family per model · frozen competence gates still FAIL
Blue = true; red = false. Crosses = wrong original label. No calibration threshold fitted; no causal intervention executed.