Audit what the measurement can establish before tracing a mechanism

Saved Qwen/Llama scores: calibration competitors and question-removal control. New Llama screen: 0 causal interventions.

A Readout changes measured success

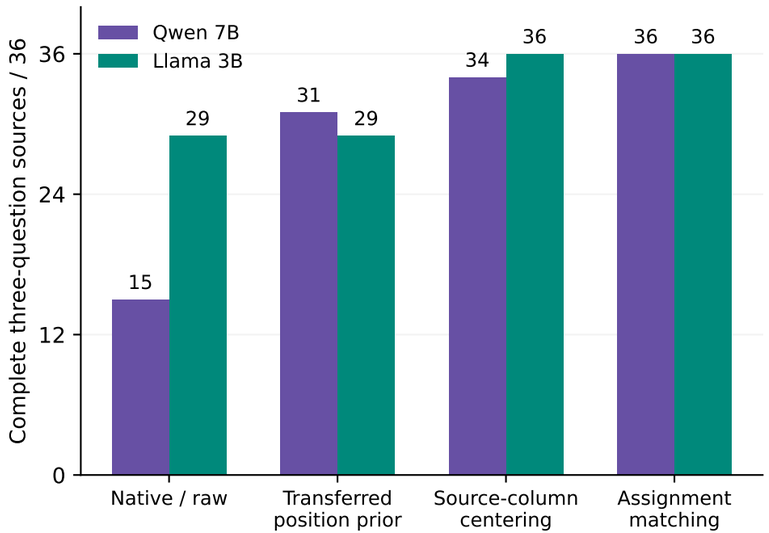

B A missing query can be filled in

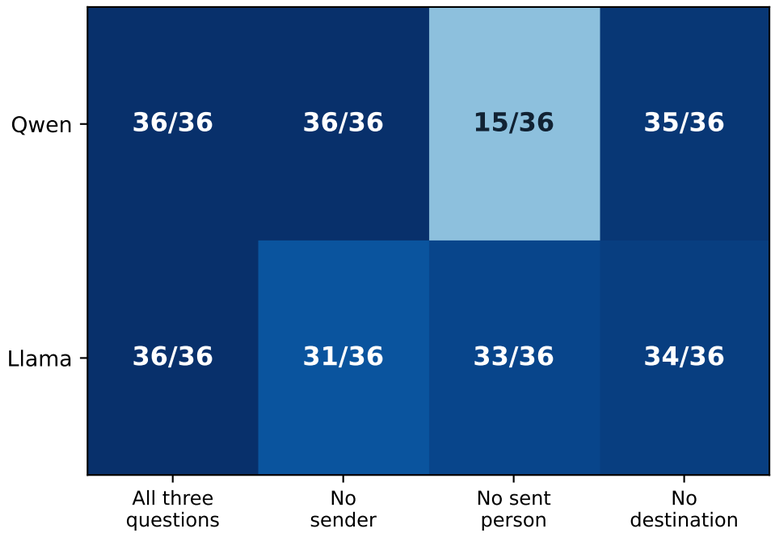

C New pilot fails native validation

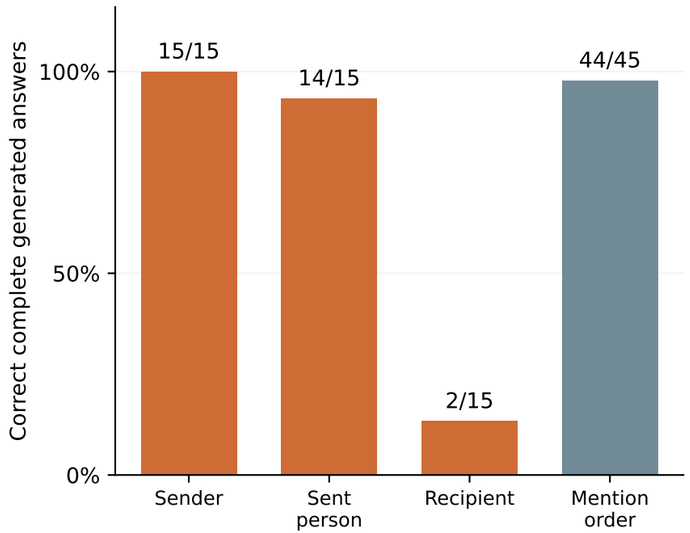

A–B: exploratory analyses of the same exposed 36 sources per checkpoint; corrections are external decisions, not new generated answers.

C: 75/90 observed answers, including one recovered response; fixed-order incomplete panel. The frozen competence gate is impossible to pass.

Next assay: four candidate people, three queried roles; change one target while the other two answers stay fixed. Design only.